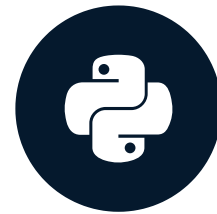


# A tale of two variables

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# Swedish motor insurance data

- Each row represents one geographic region in Sweden.
- There are 63 rows.

n_claims	total_payment_sek
108	392.5
19	46.2
13	15.7
124	422.2
40	119.4
...	...

# Descriptive statistics

```
import pandas as pd  
print(swedish_motor_insurance.mean())
```

```
n_claims          22.904762  
total_payment_sek  98.187302  
dtype: float64
```

```
print(swedish_motor_insurance['n_claims'].corr(swedish_motor_insurance['total_payment_sek']))
```

```
0.9128782350234068
```

# What is regression?

- Statistical models to explore the relationship a response variable and some explanatory variables.
- Given values of explanatory variables, you can predict the values of the response variable.

n_claims	total_payment_sek
108	3925
19	462
13	157
124	4222
40	1194
200	???

# Jargon

## **Response variable (a.k.a. dependent variable)**

The variable that you want to predict.

## **Explanatory variables (a.k.a. independent variables)**

The variables that explain how the response variable will change.

# Linear regression and logistic regression

## Linear regression

- The response variable is numeric.

## Logistic regression

- The response variable is logical.

## Simple linear/logistic regression

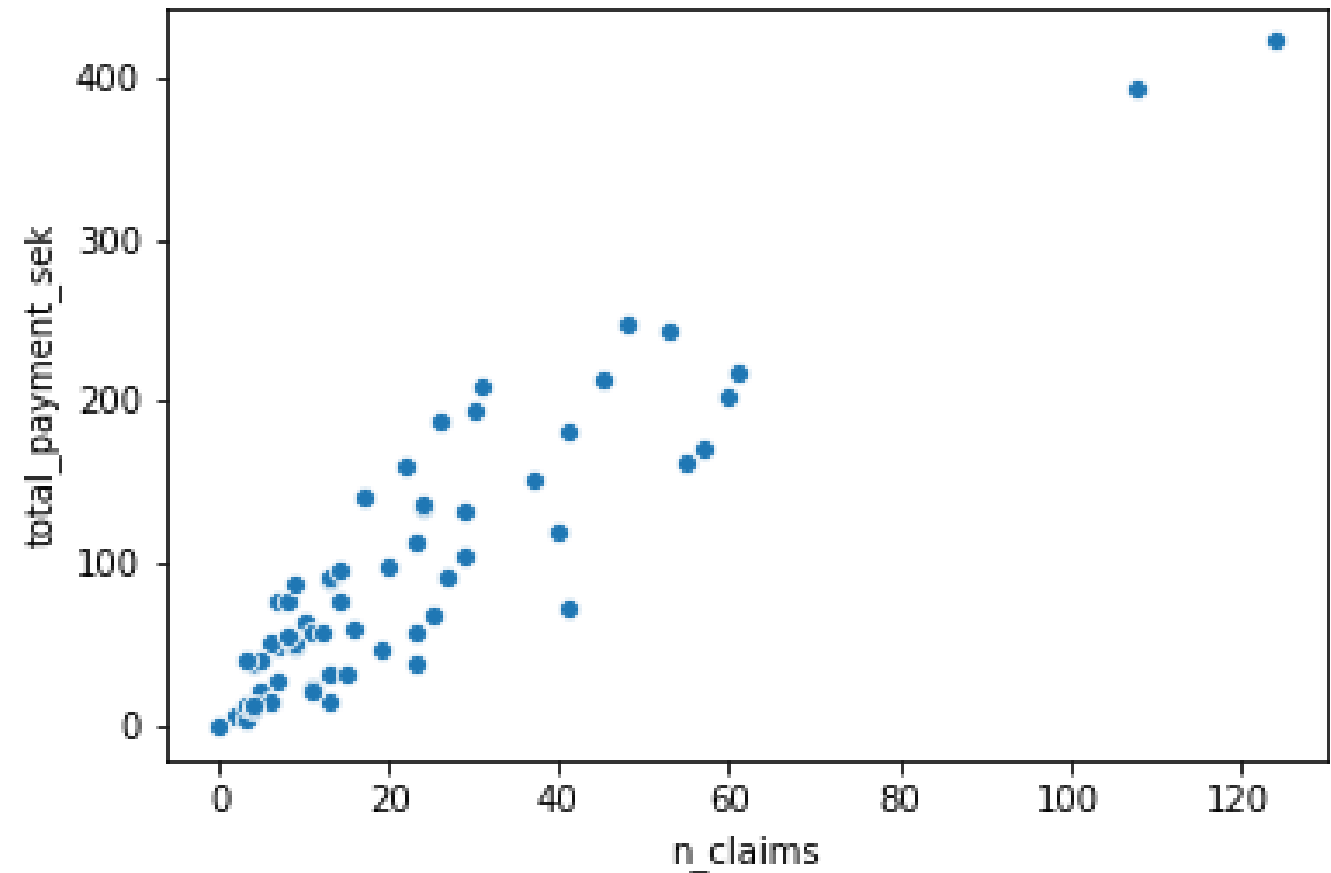
- There is only one explanatory variable.

# Visualizing pairs of variables

```
import matplotlib.pyplot as plt
import seaborn as sns

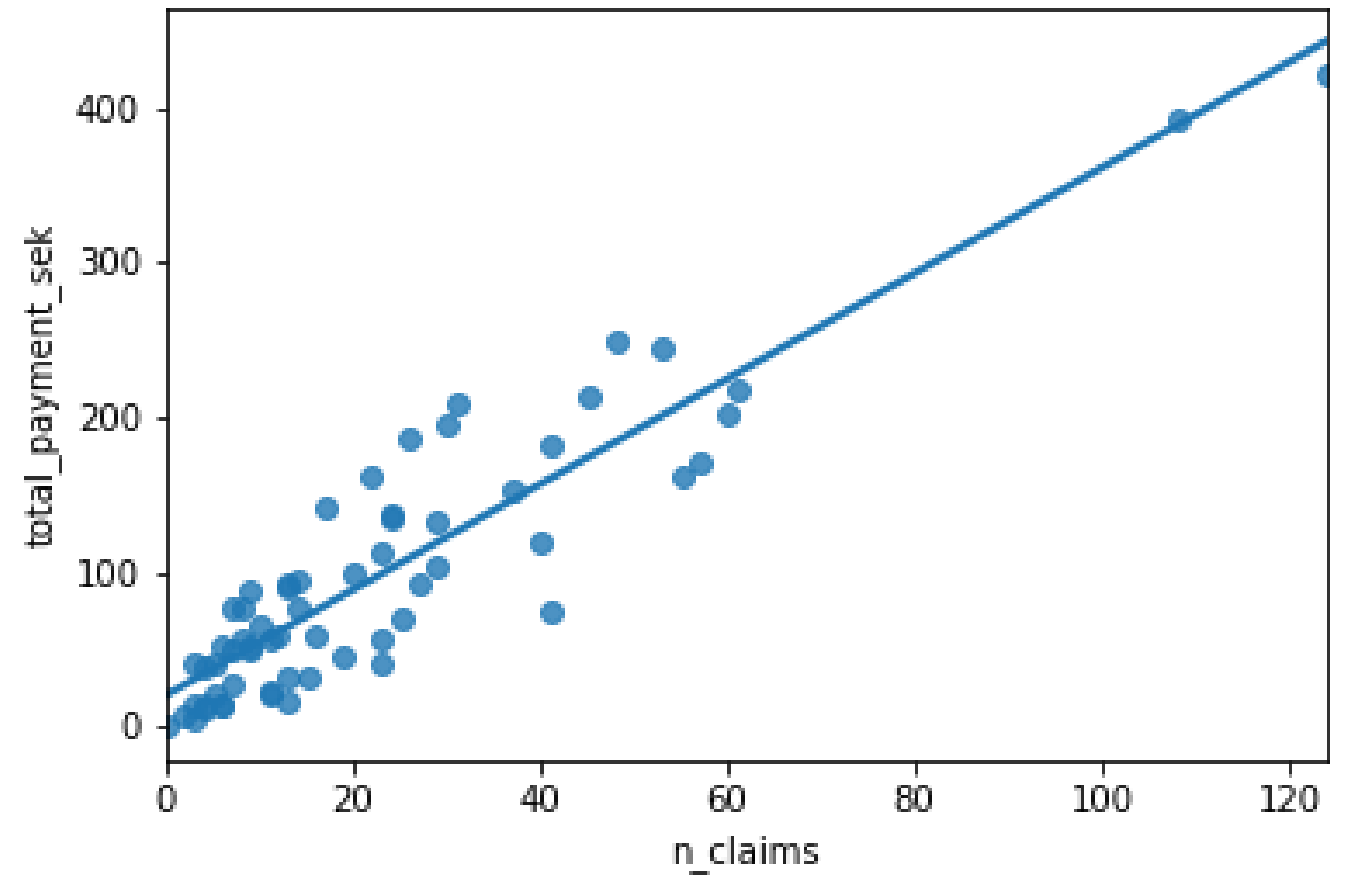
sns.scatterplot(x="n_claims",
                y="total_payment_sek",
                data=swedish_motor_insurance)

plt.show()
```



# Adding a linear trend line

```
sns.regplot(x="n_claims",  
            y="total_payment_sek",  
            data=swedish_motor_insurance,  
            ci=None)
```





# Course flow

## Chapter 1

Visualizing and fitting linear regression models.

## Chapter 2

Making predictions from linear regression models and understanding model coefficients.

## Chapter 3

Assessing the quality of the linear regression model.

## Chapter 4

Same again, but with logistic regression models

# Python packages for regression

## `statsmodels`

- Optimized for insight (focus in this course)

## `scikit-learn`

- Optimized for prediction (focus in other DataCamp courses)

# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Fitting a linear regression

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# Straight lines are defined by two things

## Intercept

The  $y$  value at the point when  $x$  is zero.


## Slope

The amount the  $y$  value increases if you increase  $x$  by one.

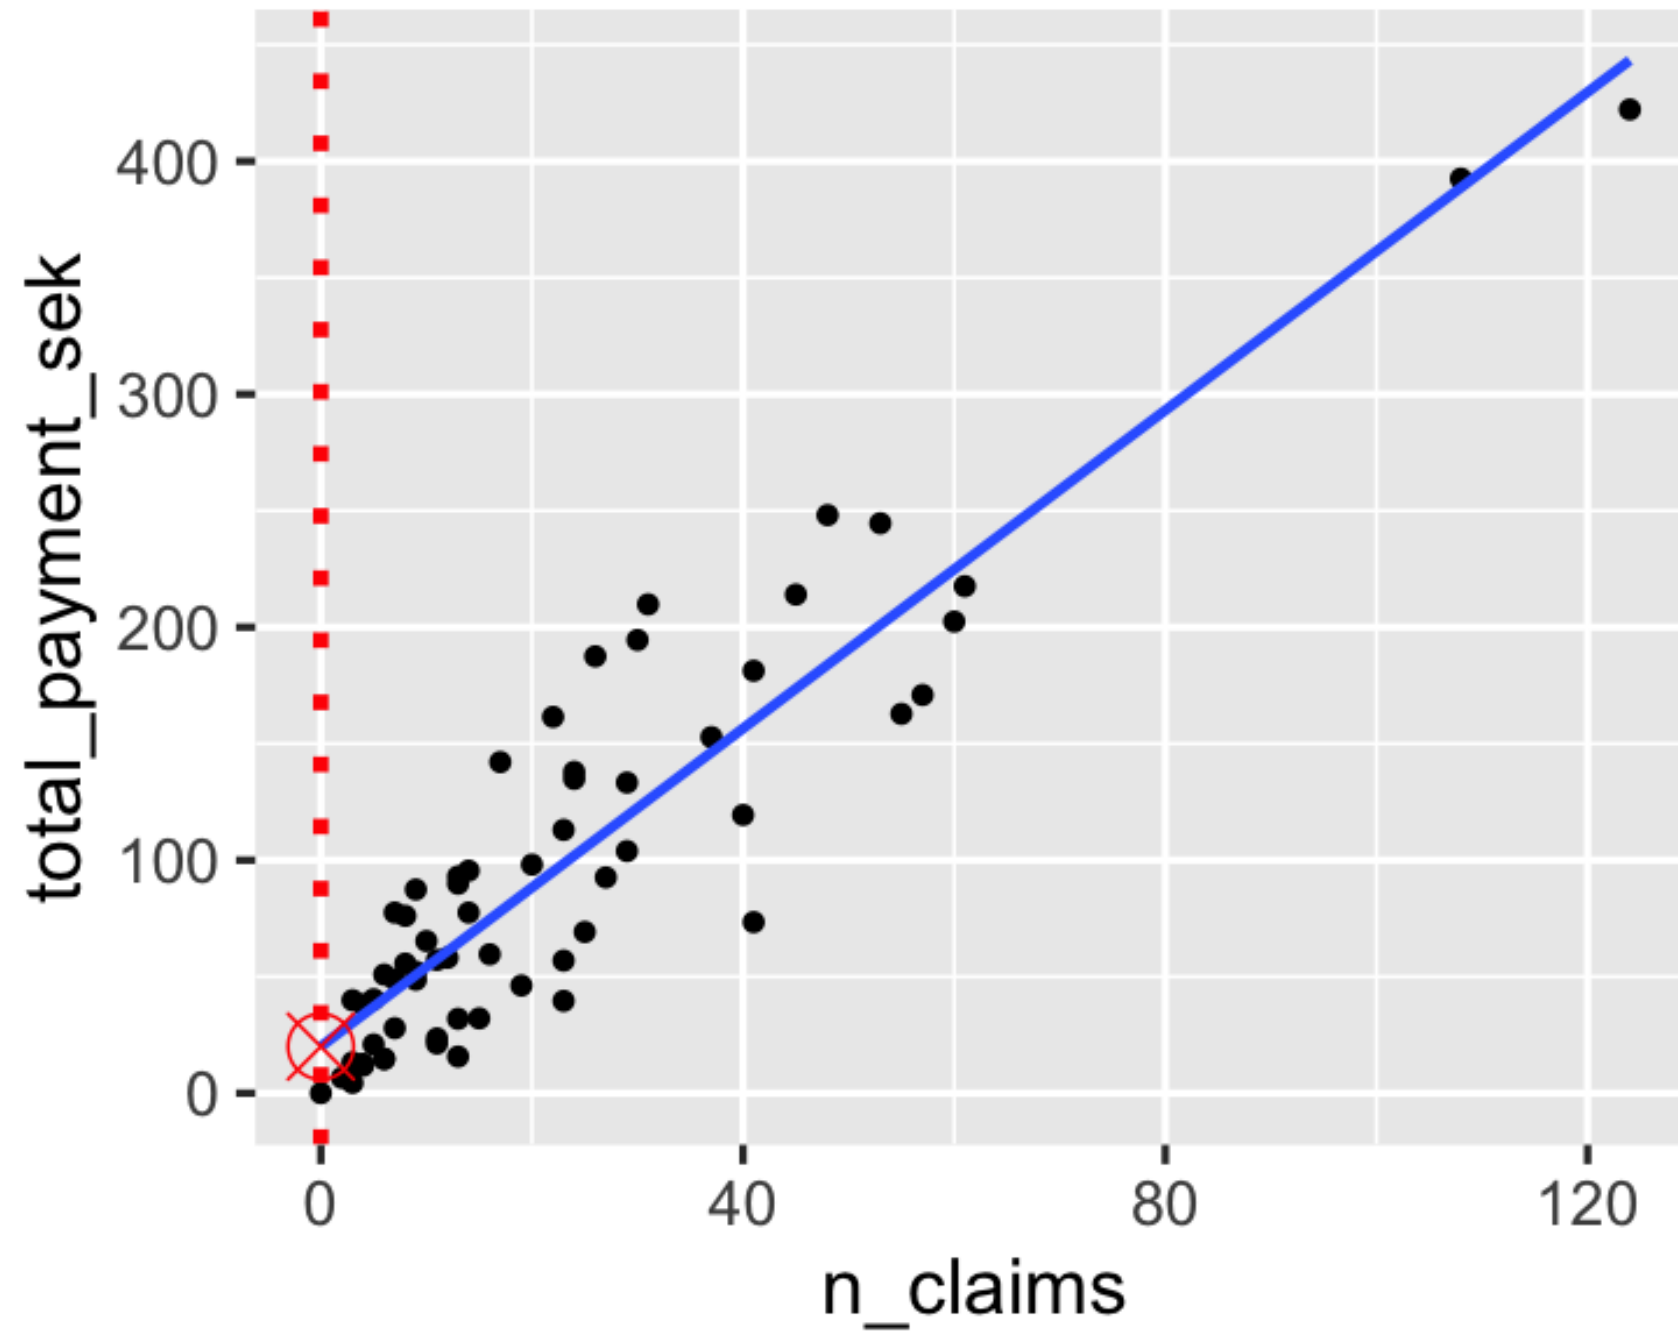
## Equation

$$y = \text{intercept} + \text{slope} * x$$

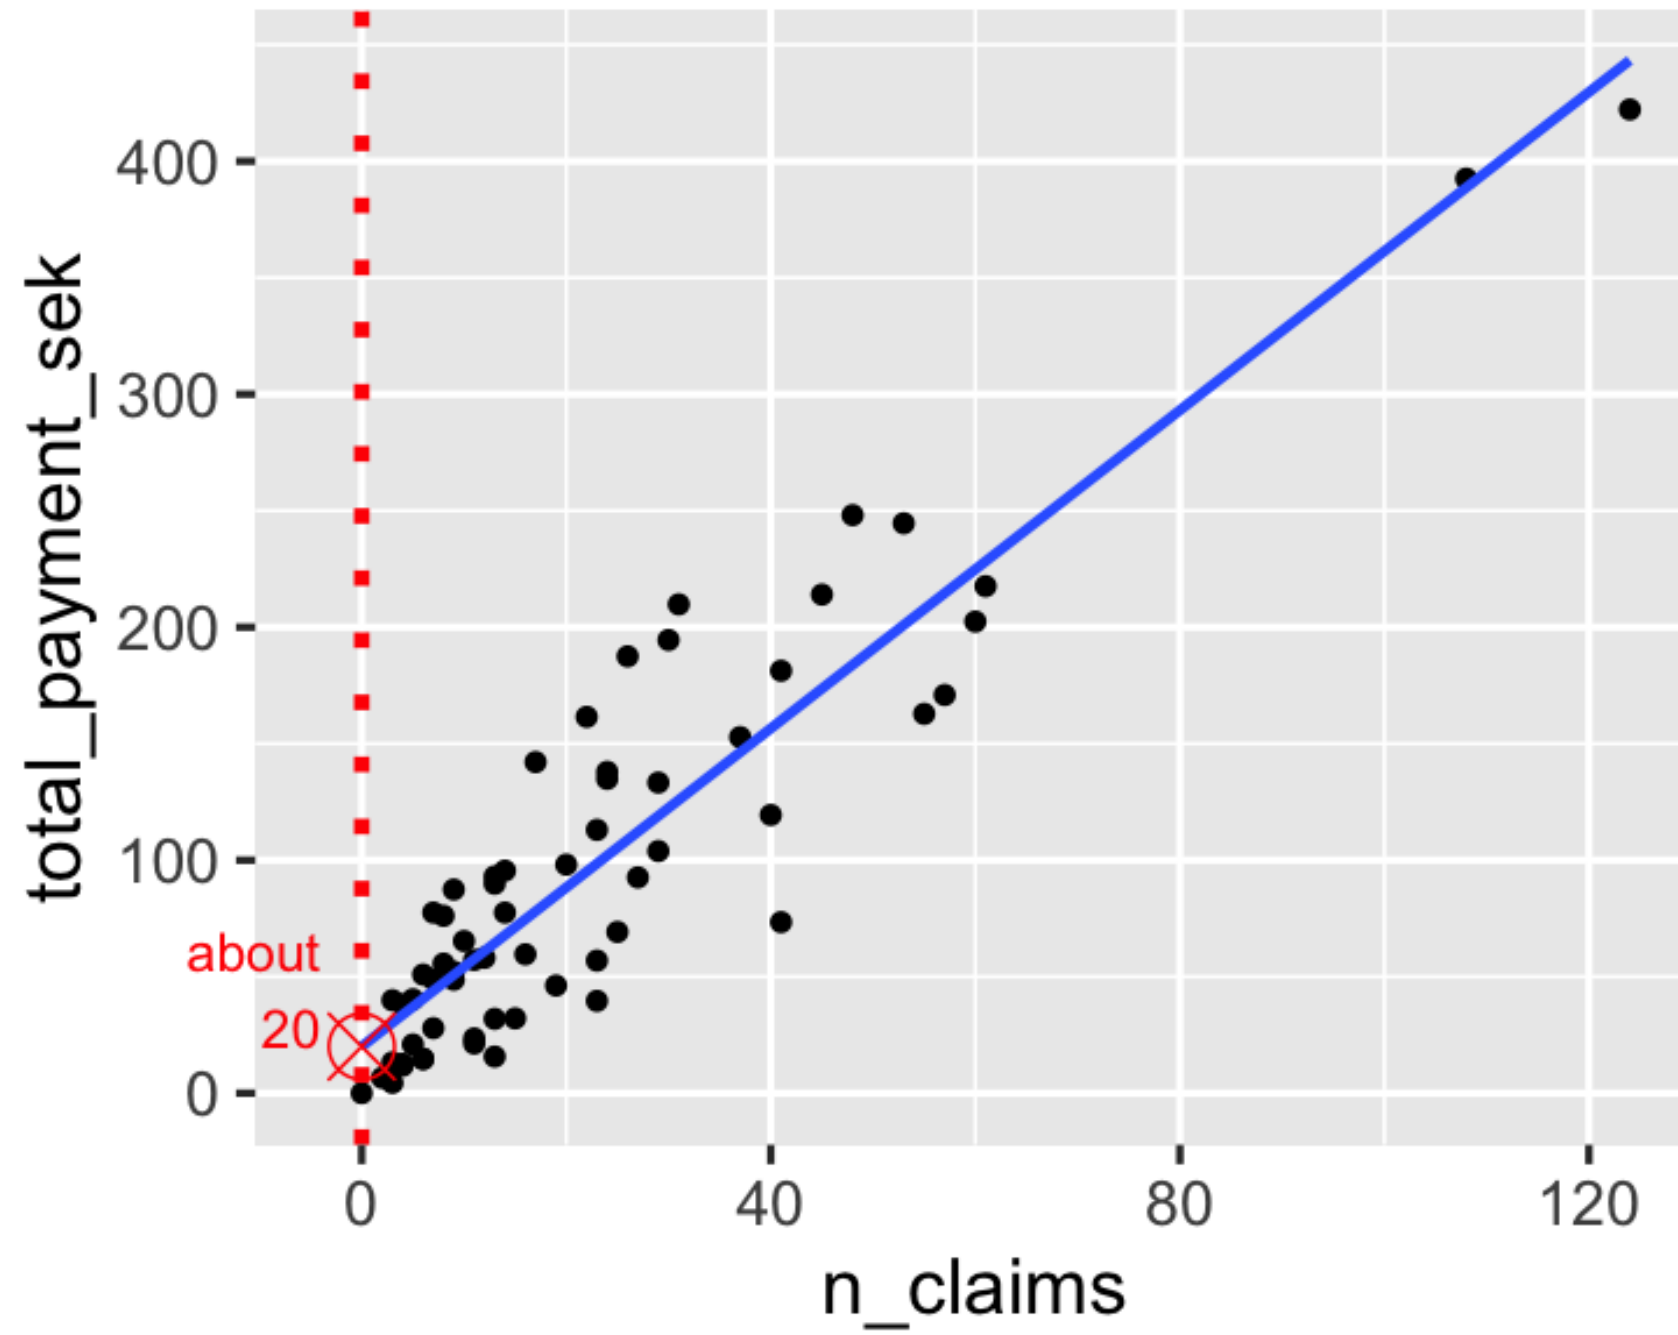
# Estimating the intercept



# Estimating the intercept

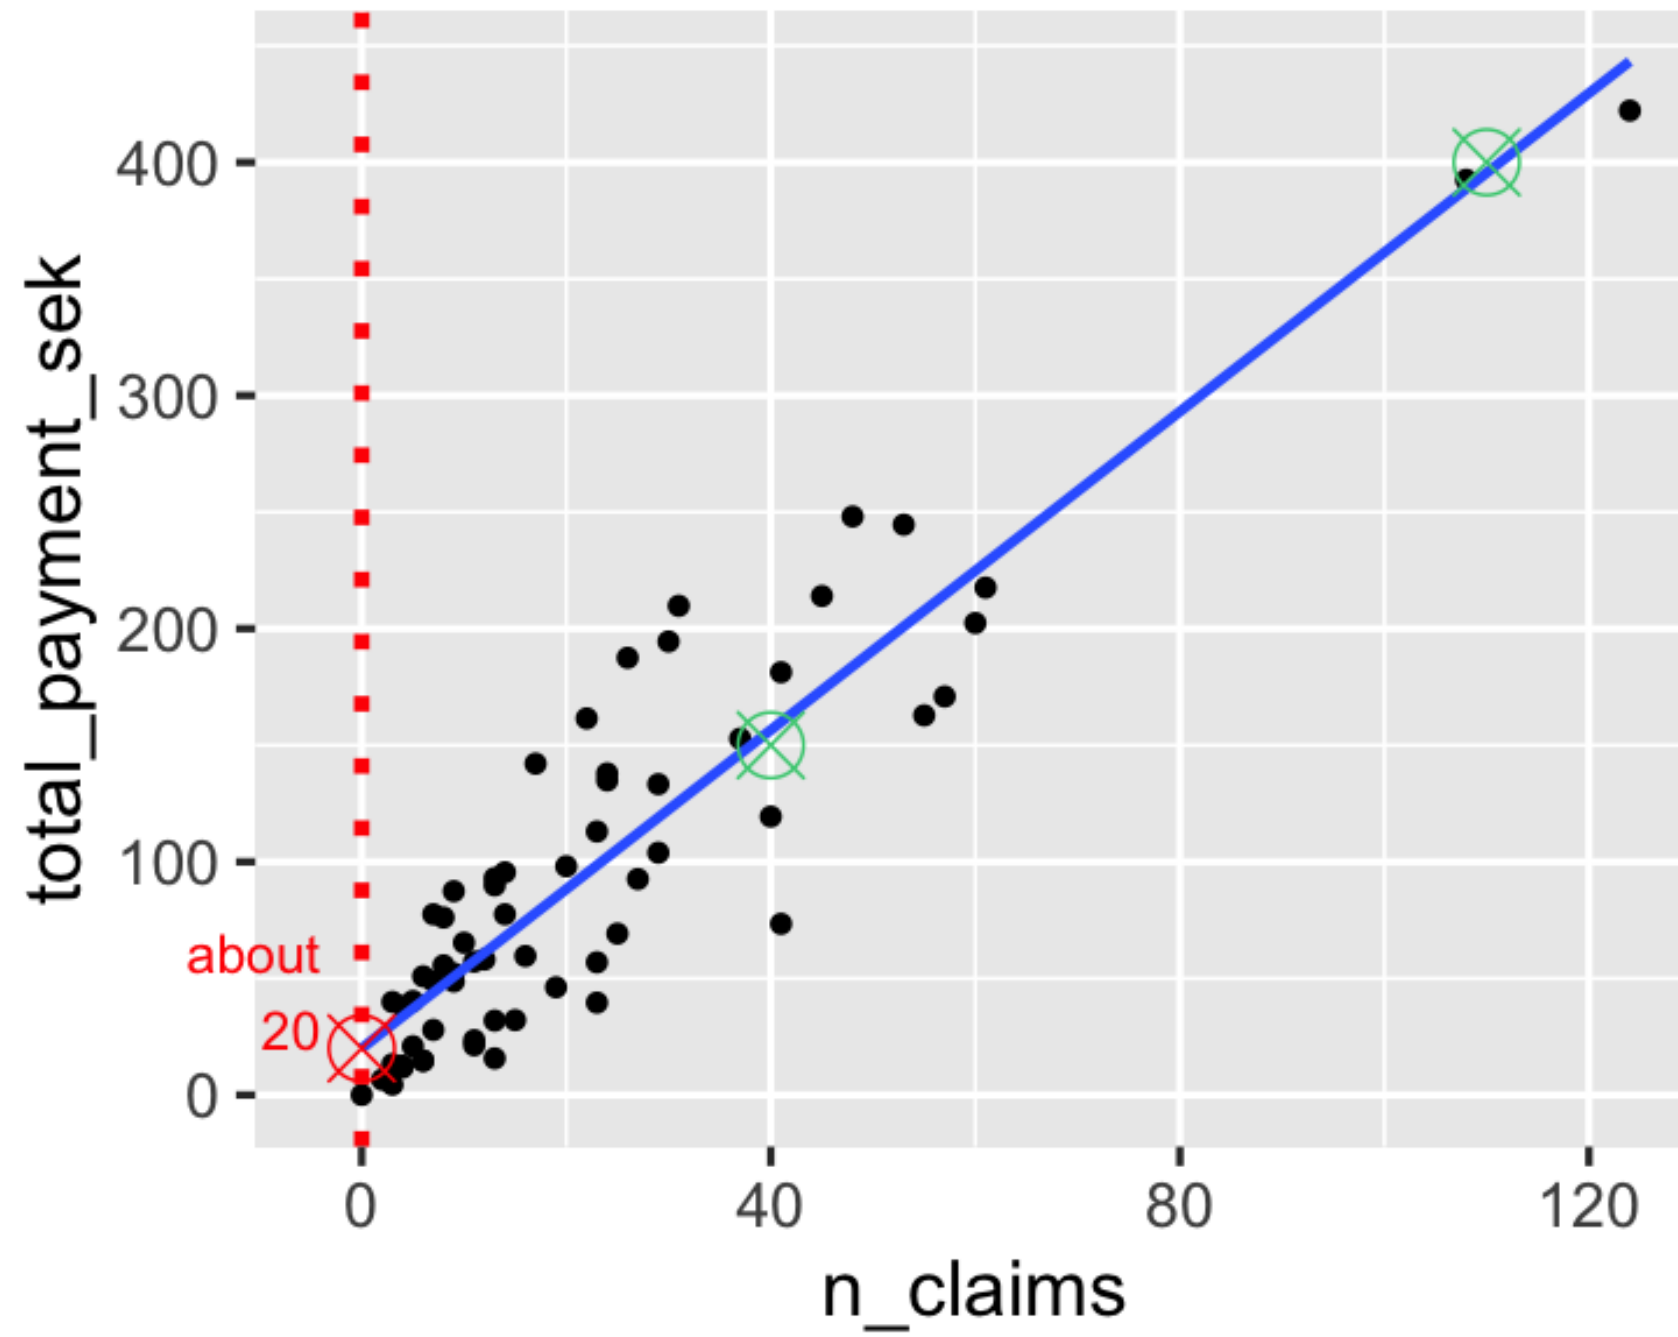


# Estimating the intercept

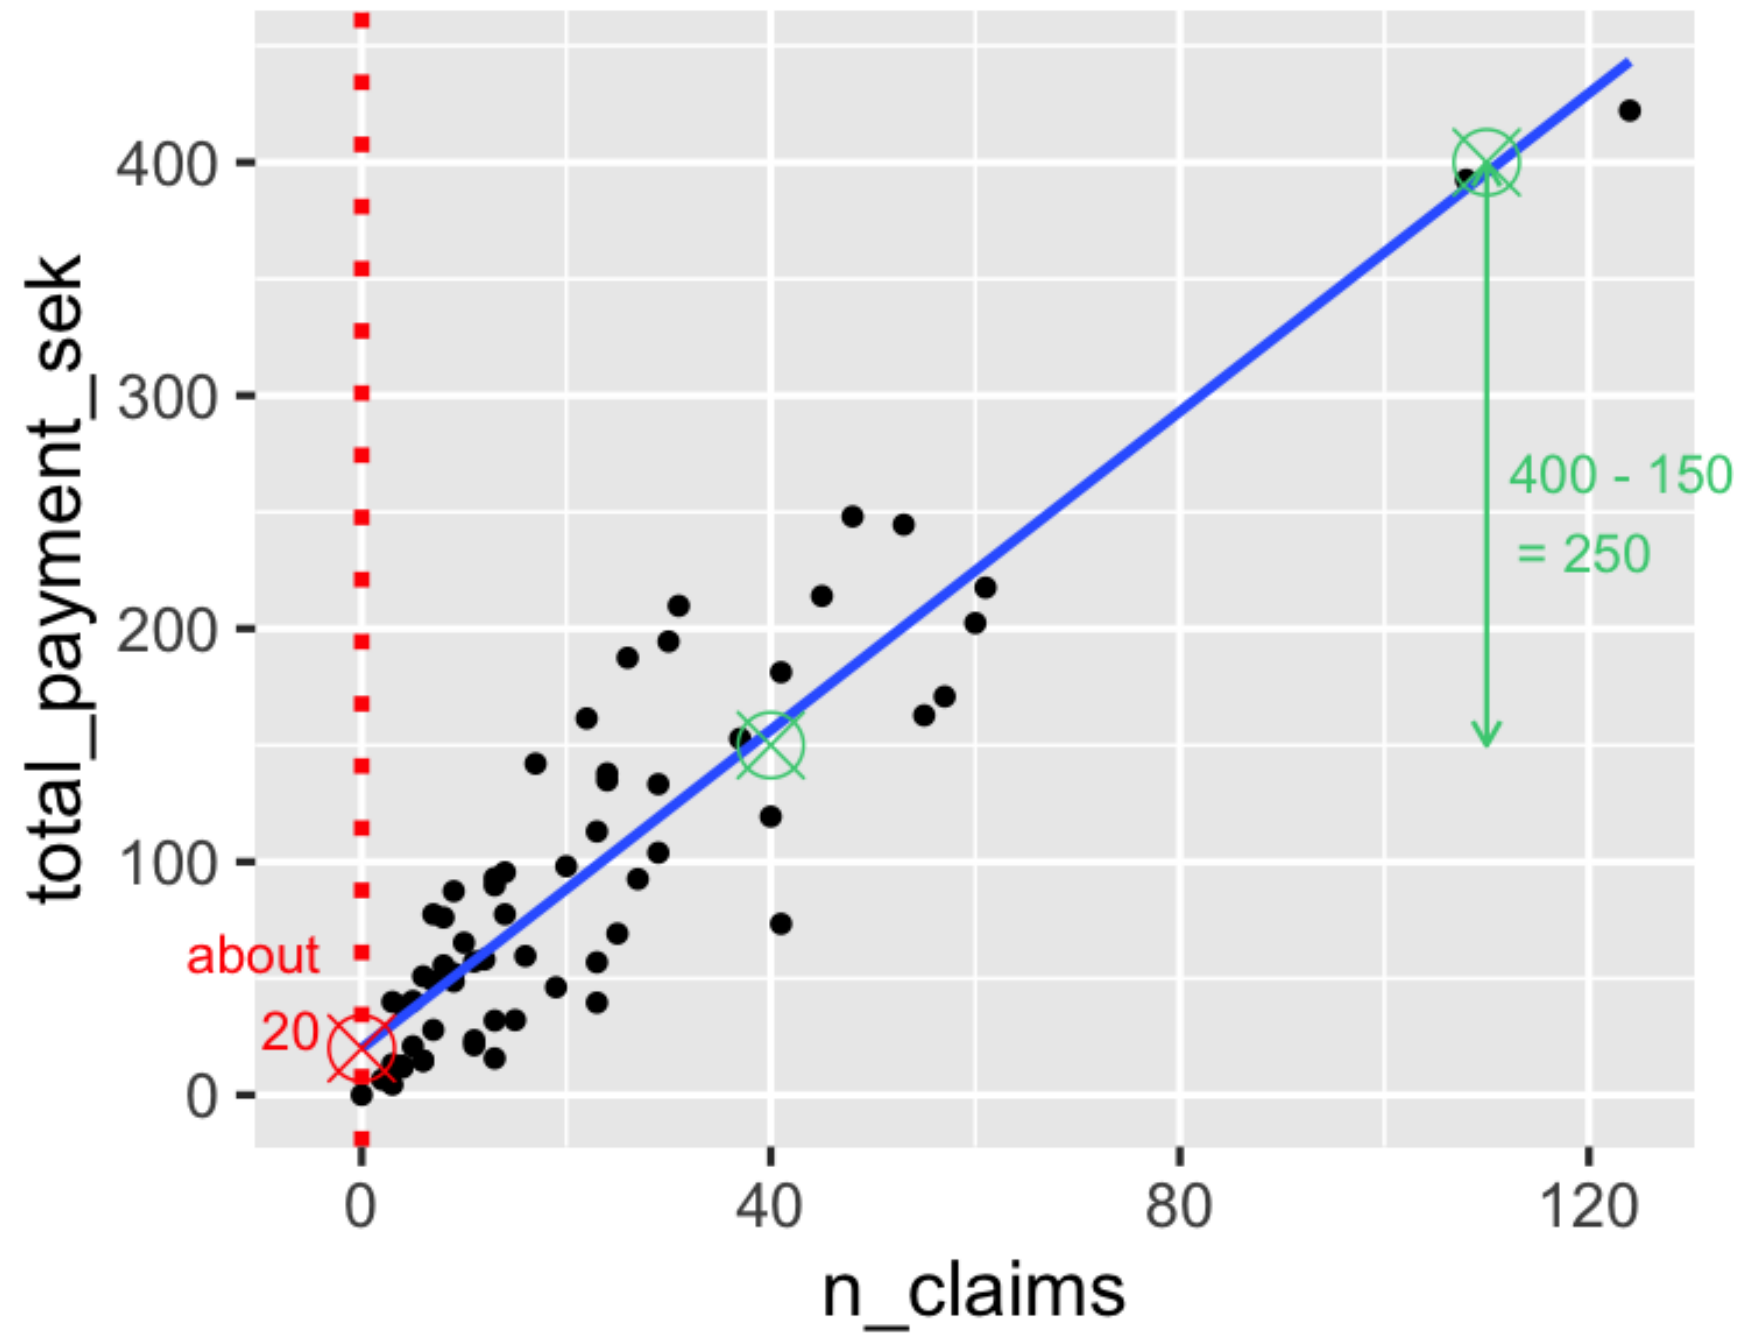




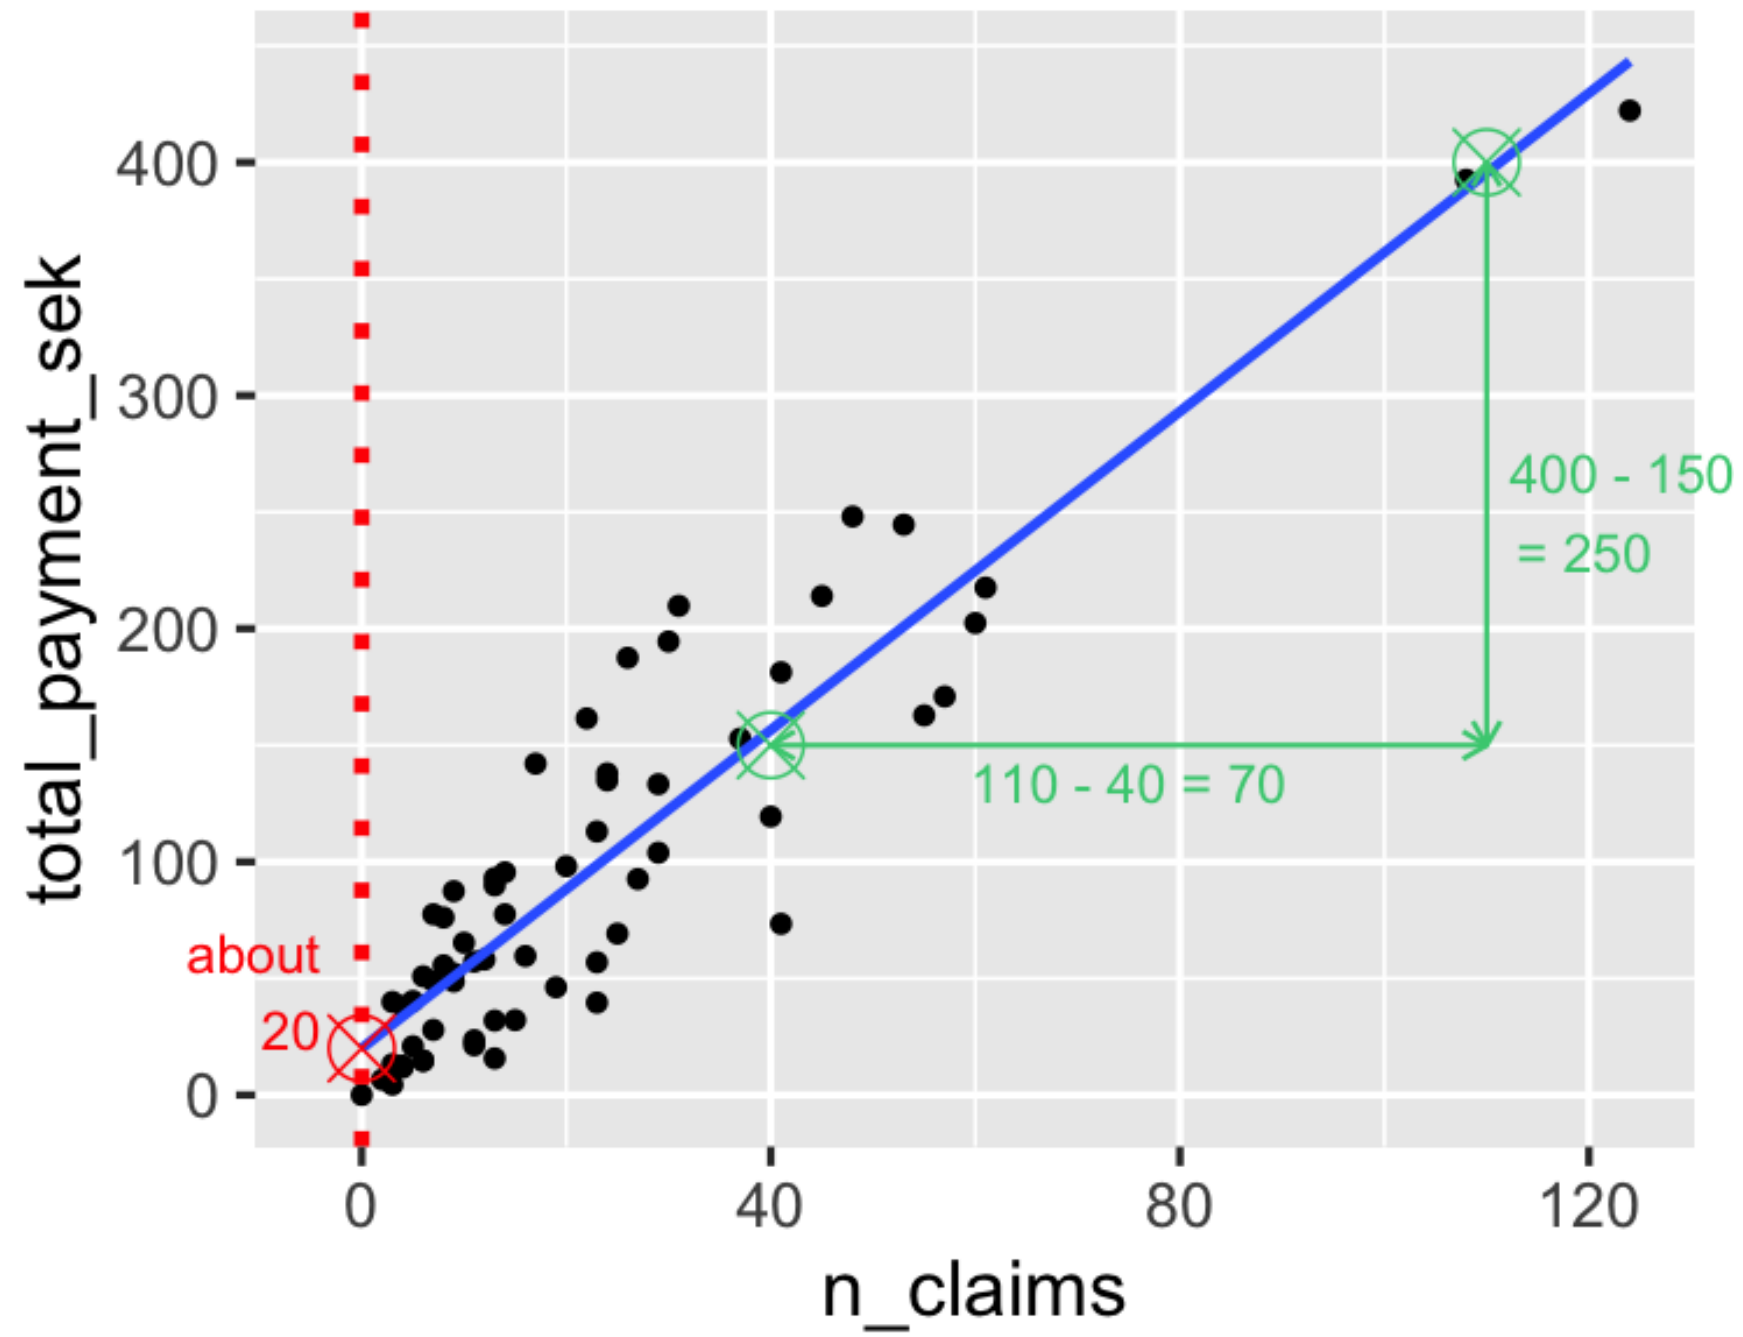
# Estimating the slope




# Estimating the slope



# Estimating the slope



# Estimating the slope



# Running a model

```
from statsmodels.formula.api import ols
mdl_payment_vs_claims = ols("total_payment_sek ~ n_claims",
                             data=swedish_motor_insurance)

mdl_payment_vs_claims = mdl_payment_vs_claims.fit()
print(mdl_payment_vs_claims.params)
```

```
Intercept    19.994486
n_claims      3.413824
dtype: float64
```

# Interpreting the model coefficients

```
Intercept    19.994486  
n_claims     3.413824  
dtype: float64
```

## Equation

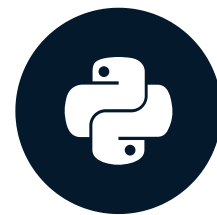
$$\text{total\_payment\_sek} = 19.99 + 3.41 * \text{n\_claims}$$

# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Categorical explanatory variables

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp



# Fish dataset

- Each row represents one fish.
- There are 128 rows in the dataset.
- There are 4 species of fish:
  - Common Bream
  - European Perch
  - Northern Pike
  - Common Roach

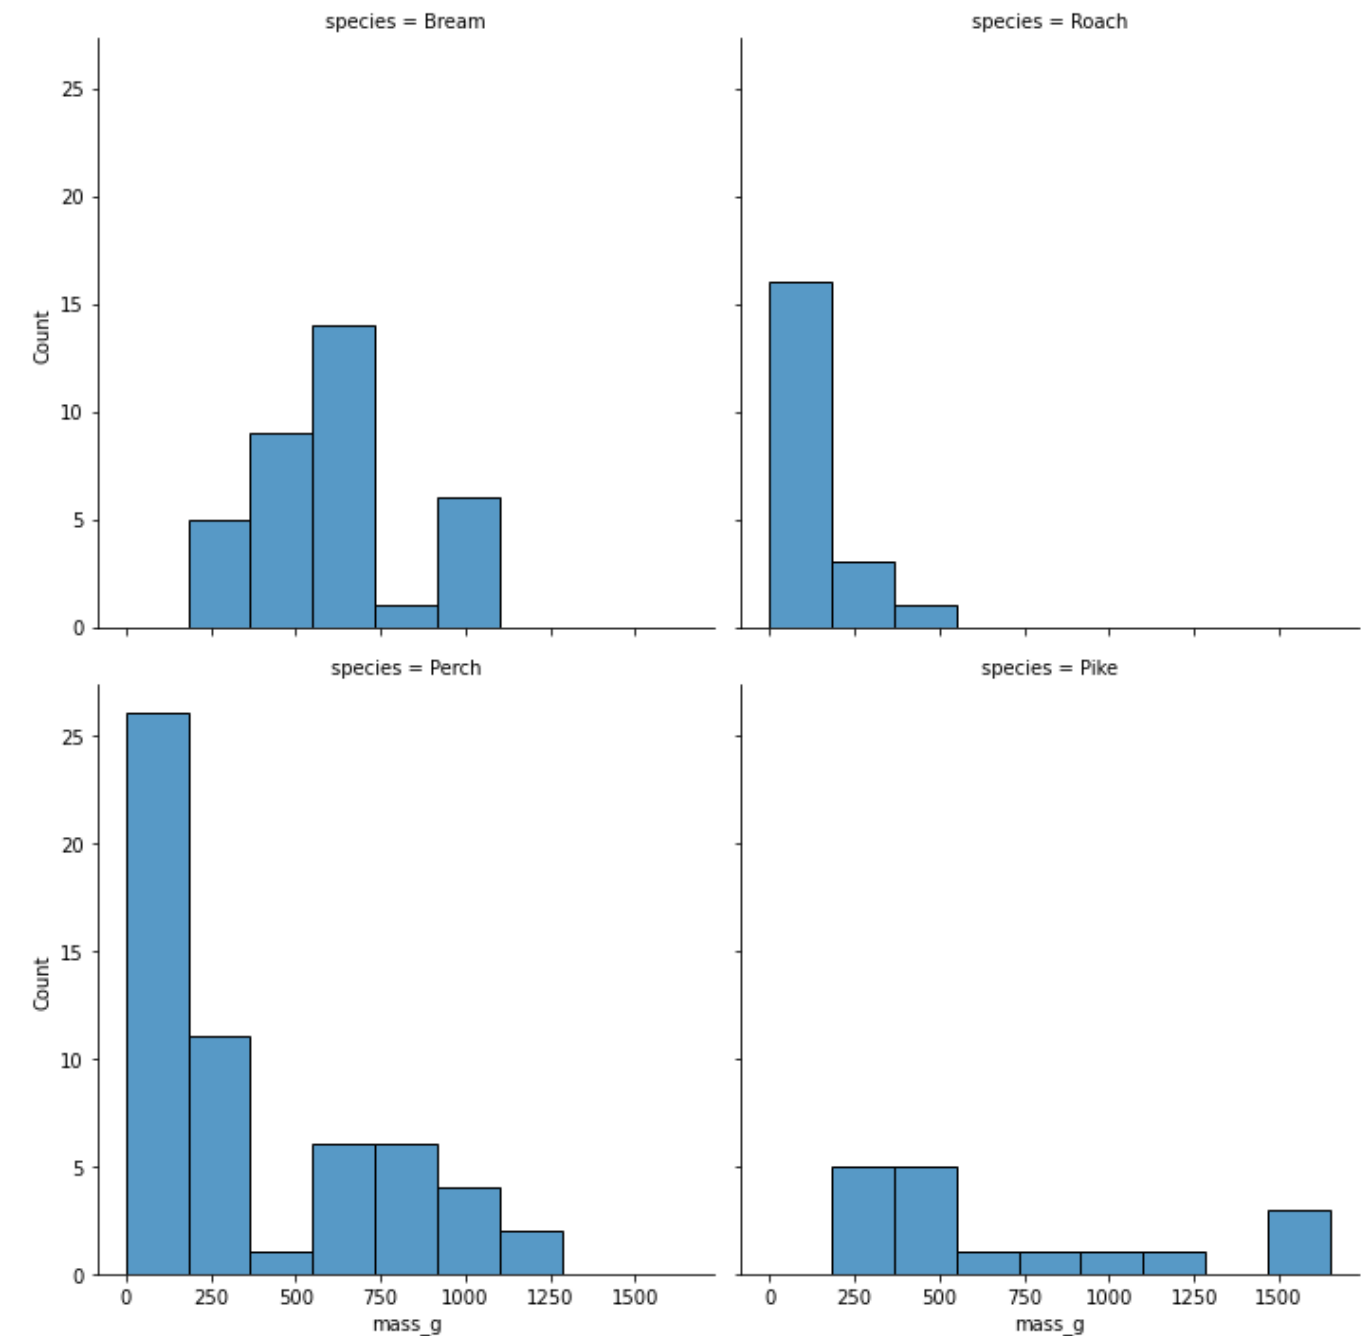
species	mass_g
Bream	242.0
Perch	5.9
Pike	200.0
Roach	40.0
...	...

# Visualizing 1 numeric and 1 categorical variable

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
sns.displot(data=fish,
            x="mass_g",
            col="species",
            col_wrap=2,
            bins=9)
```

```
plt.show()
```



# Summary statistics: mean mass by species

```
summary_stats = fish.groupby("species")["mass_g"].mean()  
print(summary_stats)
```

```
species  
Bream      617.828571  
Perch      382.239286  
Pike       718.705882  
Roach      152.050000  
Name: mass_g, dtype: float64
```

# Linear regression

```
from statsmodels.formula.api import ols
mdl_mass_vs_species = ols("mass_g ~ species", data=fish).fit()
print(mdl_mass_vs_species.params)
```

```
Intercept          617.828571
species[T.Perch]   -235.589286
species[T.Pike]    100.877311
species[T.Roach]   -465.778571
```

# Model with or without an intercept

From previous slide, model with intercept

```
mdl_mass_vs_species = ols(  
    "mass_g ~ species", data=fish).fit()  
print(mdl_mass_vs_species.params)
```

Intercept	617.828571
species[T.Perch]	-235.589286
species[T.Pike]	100.877311
species[T.Roach]	-465.778571

The coefficients are relative to the intercept:  
 $617.83 - 235.59 = 382.24!$

Model without an intercept

```
mdl_mass_vs_species = ols(  
    "mass_g ~ species + 0", data=fish).fit()  
print(mdl_mass_vs_species.params)
```

species[Bream]	617.828571
species[Perch]	382.239286
species[Pike]	718.705882
species[Roach]	152.050000

In case of a single, categorical variable,  
coefficients are the means.

# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Making predictions

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# The fish dataset: bream

```
bream = fish[fish["species"] == "Bream"]  
print(bream.head())
```


	species	mass_g	length_cm
0	Bream	242.0	23.2
1	Bream	290.0	24.0
2	Bream	340.0	23.9
3	Bream	363.0	26.3
4	Bream	430.0	26.5





# Plotting mass vs. length

```
sns.regplot(x="length_cm",  
            y="mass_g",  
            data=bream,  
            ci=None)  
  
plt.show()
```



# Running the model

```
mdl_mass_vs_length = ols("mass_g ~ length_cm", data=bream).fit()  
print(mdl_mass_vs_length.params)
```

```
Intercept    -1035.347565  
length_cm      54.549981  
dtype: float64
```

# Data on explanatory values to predict

If I set the explanatory variables to these values,  
what value would the response variable have?

```
explanatory_data = pd.DataFrame({"length_cm": np.arange(20, 41)})
```

	length_cm
0	20
1	21
2	22
3	23
4	24
5	25
...	

# Call predict()

```
print mdl_mass_vs_length.predict(explanatory_data))
```

```
0      55.652054
1     110.202035
2     164.752015
3     219.301996
4     273.851977
...
16     928.451749
17     983.001730
18    1037.551710
19    1092.101691
20    1146.651672
Length: 21, dtype: float64
```

# Predicting inside a DataFrame


```
explanatory_data = pd.DataFrame(  
    {"length_cm": np.arange(20, 41)}  
)  
prediction_data = explanatory_data.assign(  
    mass_g=mdl_mass_vs_length.predict(explanatory_data)  
)  
print(prediction_data)
```

	length_cm	mass_g
0	20	55.652054
1	21	110.202035
2	22	164.752015
3	23	219.301996
4	24	273.851977
..	...	...
16	36	928.451749
17	37	983.001730
18	38	1037.551710
19	39	1092.101691
20	40	1146.651672

# Showing predictions

```
import matplotlib.pyplot as plt
import seaborn as sns
fig = plt.figure()
sns.regplot(x="length_cm",
            y="mass_g",
            ci=None,
            data=bream,)
sns.scatterplot(x="length_cm",
               y="mass_g",
               data=prediction_data,
               color="red",
               marker="s")

plt.show()
```



# Extrapolating


*Extrapolating* means making predictions outside the range of observed data.

```
little_bream = pd.DataFrame({"length_cm": [10]})

pred_little_bream = little_bream.assign(
    mass_g=mdl_mass_vs_length.predict(little_bream))

print(pred_little_bream)
```

	length_cm	mass_g
0	10	-489.847756



# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



# Working with model objects

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# .params attribute

```
from statsmodels.formula.api import ols
mdl_mass_vs_length = ols("mass_g ~ length_cm", data = bream).fit()
print(mdl_mass_vs_length.params)
```

```
Intercept    -1035.347565
length_cm      54.549981
dtype: float64
```

# .fittedvalues attribute

*Fitted values:* predictions on the original dataset

```
print mdl_mass_vs_length.fittedvalues
```

or equivalently

```
explanatory_data = bream["length_cm"]  
  
print(mdl_mass_vs_length.predict(explanatory_data))
```

```
0      230.211993  
1      273.851977  
2      268.396979  
3      399.316934  
4      410.226930  
  
...  
30     873.901768  
31     873.901768  
32     939.361745  
33    1004.821722  
34    1037.551710  
Length: 35, dtype: float64
```

# .resid attribute


*Residuals*: actual response values minus predicted response values

```
print mdl_mass_vs_length.resid
```

or equivalently

```
print(bream["mass_g"] - mdl_mass_vs_length.fittedvalues)
```

```
0    11.788007
1    16.148023
2    71.603021
3   -36.316934
4    19.773070
...
```



# .summary()

```
mdl_mass_vs_length.summary()
```

```

                    OLS Regression Results
=====
Dep. Variable:      mass_g      R-squared:      0.878
Model:              OLS        Adj. R-squared:    0.874
Method:             Least Squares    F-statistic: 237.6
Date:              Thu, 29 Oct 2020    Prob (F-statistic): 1.22e-16
Time:              13:23:21    Log-Likelihood: -199.35
No. Observations:   35    AIC:      402.7
Df Residuals:       33    BIC:      405.8
Df Model:           1
Covariance Type:    nonrobust
=====
                    coef    std err          t      P>|t|      [0.025     0.975]
-----
Intercept  -1035.3476    107.973     -9.589     0.000   -1255.020    -815.676
length_cm    54.5500      3.539     15.415     0.000     47.350     61.750
=====
Omnibus:         7.314    Durbin-Watson:      1.478
Prob(Omnibus):    0.026    Jarque-Bera (JB):    10.857
Skew:            -0.252    Prob(JB):           0.00439
Kurtosis:         5.682    Cond. No.           263.
```

## OLS Regression Results

```
=====
Dep. Variable:          mass_g    R-squared:                0.878
Model:                  OLS       Adj. R-squared:           0.874
Method:                 Least Squares    F-statistic:            237.6
Date:                   Thu, 29 Oct 2020    Prob (F-statistic):      1.22e-16
Time:                   13:23:21    Log-Likelihood:         -199.35
No. Observations:       35         AIC:                     402.7
Df Residuals:           33         BIC:                     405.8
Df Model:               1
Covariance Type:        nonrobust
```

	coef	std err	t	P> t	[0.025	0.975]
<-----						
Intercept	-1035.3476	107.973	-9.589	0.000	-1255.020	-815.676
length_cm	54.5500	3.539	15.415	0.000	47.350	61.750
=====						
Omnibus:		7.314	Durbin-Watson:			1.478
Prob(Omnibus):		0.026	Jarque-Bera (JB):			10.857
Skew:		-0.252	Prob(JB):			0.00439
Kurtosis:		5.682	Cond. No.			263.


# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



# Regression to the mean

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# The concept

- Response value = fitted value + residual
- "The stuff you explained" + "the stuff you couldn't explain"
- Residuals exist due to problems in the model *and* fundamental randomness
- Extreme cases are often due to randomness
- *Regression to the mean* means extreme cases don't persist over time

# Pearson's father son dataset

- 1078 father/son pairs
- Do tall fathers have tall sons?

father_height_cm	son_height_cm
165.2	151.8
160.7	160.6
165.0	160.9
167.0	159.5
155.3	163.3
...	...

<sup>1</sup> Adapted from <https://www.rdocumentation.org/packages/UsingR/topics/father.son>


# Scatter plot

```
fig = plt.figure()
```

```
sns.scatterplot(x="father_height_cm",  
                y="son_height_cm",  
                data=father_son)
```

```
plt.axline(xy1=(150, 150),  
           slope=1,  
           linewidth=2,  
           color="green")
```

```
plt.axis("equal")  
plt.show()
```




# Adding a regression line

```
fig = plt.figure()

sns.regplot(x="father_height_cm",
            y="son_height_cm",
            data=father_son,
            ci = None,
            line_kws={"color": "black"})

plt.axline(xy1 = (150, 150),
           slope=1,
           linewidth=2,
           color="green")

plt.axis("equal")
plt.show()
```



# Running a regression

```
mdl_son_vs_father = ols("son_height_cm ~ father_height_cm",  
                        data = father_son).fit()  
  
print(mdl_son_vs_father.params)
```

```
Intercept          86.071975  
father_height_cm    0.514093  
dtype: float64
```

# Making predictions

```
really_tall_father = pd.DataFrame(  
    {"father_height_cm": [190]})
```

```
mdl_son_vs_father.predict(  
    really_tall_father)
```

183.7

```
really_short_father = pd.DataFrame(  
    {"father_height_cm": [150]})
```

```
mdl_son_vs_father.predict(  
    really_short_father)
```

163.2

# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



# Transforming variables

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# Perch dataset


```
perch = fish[fish["species"] == "Perch"]  
print(perch.head())
```

	species	mass_g	length_cm
55	Perch	5.9	7.5
56	Perch	32.0	12.5
57	Perch	40.0	13.8
58	Perch	51.5	15.0
59	Perch	70.0	15.7



# It's not a linear relationship

```
sns.regplot(x="length_cm",  
            y="mass_g",  
            data=perch,  
            ci=None)  
  
plt.show()
```



# Bream vs. perch






# Plotting mass vs. length cubed

```
perch["length_cm_cubed"] = perch["length_cm"] ** 3
```

```
sns.regplot(x="length_cm_cubed",  
            y="mass_g",  
            data=perch,  
            ci=None)  
  
plt.show()
```



# Modeling mass vs. length cubed

```
perch["length_cm_cubed"] = perch["length_cm"] ** 3

mdl_perch = ols("mass_g ~ length_cm_cubed", data=perch).fit()
mdl_perch.params
```

```
Intercept          -0.117478
length_cm_cubed     0.016796
dtype: float64
```

# Predicting mass vs. length cubed


```
explanatory_data = pd.DataFrame({"length_cm_cubed": np.arange(10, 41, 5) ** 3,  
                                "length_cm": np.arange(10, 41, 5)})
```

```
prediction_data = explanatory_data.assign(  
    mass_g=mdl_perch.predict(explanatory_data))  
print(prediction_data)
```


	length_cm_cubed	length_cm	mass_g
0	1000	10	16.678135
1	3375	15	56.567717
2	8000	20	134.247429
3	15625	25	262.313982
4	27000	30	453.364084
5	42875	35	719.994447
6	64000	40	1074.801781

# Plotting mass vs. length cubed

```
fig = plt.figure()
sns.regplot(x="length_cm_cubed", y="mass_g",
            data=perch, ci=None)
sns.scatterplot(data=prediction_data,
                x="length_cm_cubed", y="mass_g",
                color="red", marker="s")
```



```
fig = plt.figure()
sns.regplot(x="length_cm", y="mass_g",
            data=perch, ci=None)
sns.scatterplot(data=prediction_data,
                x="length_cm", y="mass_g",
                color="red", marker="s")
```





# Facebook advertising dataset

## How advertising works


1. Pay Facebook to shows ads.
2. People see the ads ("impressions").
3. Some people who see it, click it.

- 936 rows
- Each row represents 1 advert

spent_usd	n_impressions	n_clicks
1.43	7350	1
1.82	17861	2
1.25	4259	1
1.29	4133	1
4.77	15615	3
...	...	...

# Plot is cramped

```
sns.regplot(x="spent_usd",  
            y="n_impressions",  
            data=ad_conversion,  
            ci=None)
```




# Square root vs square root

```
ad_conversion["sqrt_spent_usd"] = np.sqrt(
    ad_conversion["spent_usd"])

ad_conversion["sqrt_n_impressions"] = np.sqrt(
    ad_conversion["n_impressions"])

sns.regplot(x="sqrt_spent_usd",
            y="sqrt_n_impressions",
            data=ad_conversion,
            ci=None)
```



# Modeling and predicting

```
mdl_ad = ols("sqrt_n_impressions ~ sqrt_spent_usd", data=ad_conversion).fit()
```

```
explanatory_data = pd.DataFrame({"sqrt_spent_usd": np.sqrt(np.arange(0, 601, 100)),  
                                "spent_usd": np.arange(0, 601, 100)})
```

```
prediction_data = explanatory_data.assign(sqrt_n_impressions=mdl_ad.predict(explanatory_data),  
                                         n_impressions=mdl_ad.predict(explanatory_data) ** 2)  
print(prediction_data)
```

	sqrt_spent_usd	spent_usd	sqrt_n_impressions	n_impressions
0	0.000000	0	15.319713	2.346936e+02
1	10.000000	100	597.736582	3.572890e+05
2	14.142136	200	838.981547	7.038900e+05
3	17.320508	300	1024.095320	1.048771e+06
4	20.000000	400	1180.153450	1.392762e+06
5	22.360680	500	1317.643422	1.736184e+06
6	24.494897	600	1441.943858	2.079202e+06

# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Quantifying model fit

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON




**Maarten Van den Broeck**


Content Developer at DataCamp

# Bream and perch models

Bream



Perch



# Coefficient of determination

Sometimes called "r-squared" or "R-squared".

The proportion of the variance in the response variable that is predictable from the explanatory variable

- 1 means a perfect fit
- 0 means the worst possible fit



# .summary()

Look at the value titled "R-Squared"

```
mdl_bream = ols("mass_g ~ length_cm", data=bream).fit()
```

```
print(mdl_bream.summary())
```

```
# Some lines of output omitted
```

## OLS Regression Results

Dep. Variable:	mass_g	R-squared:	0.878
Model:	OLS	Adj. R-squared:	0.874
Method:	Least Squares	F-statistic:	237.6

# .rsquared attribute

```
print mdl_bream.rsquared)
```


```
0.8780627095147174
```

# It's just correlation squared

```
coeff_determination = bream["length_cm"].corr(bream["mass_g"]) ** 2  
print(coeff_determination)
```

```
0.8780627095147173
```

# Residual standard error (RSE)



- A "typical" difference between a prediction and an observed response
- It has the same unit as the response variable.
- $MSE = RSE^2$

# .mse\_resid attribute

```
mse = mdl_bream.mse_resid  
print('mse: ', mse)
```

```
mse: 5498.555084973521
```

```
rse = np.sqrt(mse)  
print("rse: ", rse)
```

```
rse: 74.15224261594197
```

# Calculating RSE: residuals squared

```
residuals_sq = mdl_bream.resid ** 2

print("residuals sq: \n", residuals_sq)
```

```
residuals sq:
0      138.957118
1      260.758635
2     5126.992578
3     1318.919660
4      390.974309

30     2125.047026
31     6576.923291
32      206.259713
33      889.335096
34     7665.302003
Length: 35, dtype: float64
```

# Calculating RSE: sum of residuals squared

```
residuals_sq = mdl_bream.resid ** 2

resid_sum_of_sq = sum(residuals_sq)

print("resid sum of sq :",
      resid_sum_of_sq)
```

```
resid sum of sq : 181452.31780412616
```

# Calculating RSE: degrees of freedom

```
residuals_sq = mdl_bream.resid ** 2

resid_sum_of_sq = sum(residuals_sq)

deg_freedom = len(bream.index) - 2

print("deg freedom: ", deg_freedom)
```

```
deg freedom: 33
```

*Degrees of freedom* equals the number of observations minus the number of model coefficients.



# Calculating RSE: square root of ratio

```
residuals_sq = mdl_bream.resid ** 2

resid_sum_of_sq = sum(residuals_sq)

deg_freedom = len(bream.index) - 2

rse = np.sqrt(resid_sum_of_sq/deg_freedom)

print("rse :", rse)
```

```
rse : 74.15224261594197
```

# Interpreting RSE

`mdl_bream` has an RSE of `74` .

The difference between predicted bream masses and observed bream masses is typically about 74g.

# Root-mean-square error (RMSE)

```
residuals_sq = mdl_bream.resid ** 2

resid_sum_of_sq = sum(residuals_sq)

deg_freedom = len(bream.index) - 2

rse = np.sqrt(resid_sum_of_sq/deg_freedom)

print("rse :", rse)
```

```
rse : 74.15224261594197
```

```
residuals_sq = mdl_bream.resid ** 2

resid_sum_of_sq = sum(residuals_sq)

n_obs = len(bream.index)

rmse = np.sqrt(resid_sum_of_sq/n_obs)

print("rmse :", rmse)
```

```
rmse : 72.00244396727619
```

# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Visualizing model fit

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp


# Residual properties of a good fit

- Residuals are normally distributed
- The mean of the residuals is zero

# Bream and perch again


Bream: the "good" model

```
mdl_bream = ols("mass_g ~ length_cm", data=bream).fit()
```




Perch: the "bad" model

```
mdl_perch = ols("mass_g ~ length_cm", data=perch).fit()
```




# Residuals vs. fitted

Bream




Perch






# Q-Q plot

Bream




Perch




# Scale-location plot

Bream




Perch




# residplot()

```
sns.residplot(x="length_cm", y="mass_g", data=bream, lowess=True)  
plt.xlabel("Fitted values")  
plt.ylabel("Residuals")
```




# qqplot()

```
from statsmodels.api import qqplot  
qqplot(data=mdl_bream.resid, fit=True, line="45")
```



# Scale-location plot

```
model_norm_residuals_bream = mdl_bream.get_influence().resid_studentized_internal
model_norm_residuals_abs_sqrt_bream = np.sqrt(np.abs(model_norm_residuals_bream))
sns.regplot(x=mdl_bream.fittedvalues, y=model_norm_residuals_abs_sqrt_bream, ci=None, lowess=True)
plt.xlabel("Fitted values")
plt.ylabel("Sqrt of abs val of stdized residuals")
```



# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Outliers, leverage, and influence

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# Roach dataset

```
roach = fish[fish['species'] == "Roach"]  
print(roach.head())
```


	species	mass_g	length_cm
35	Roach	40.0	12.9
36	Roach	69.0	16.5
37	Roach	78.0	17.5
38	Roach	87.0	18.2
39	Roach	120.0	18.6





# Which points are outliers?

```
sns.regplot(x="length_cm",  
            y="mass_g",  
            data=roach,  
            ci=None)  
  
plt.show()
```




# Extreme explanatory values

```
roach["extreme_l"] = ((roach["length_cm"] < 15) |  
                     (roach["length_cm"] > 26))
```

```
fig = plt.figure()  
sns.regplot(x="length_cm",  
            y="mass_g",  
            data=roach,  
            ci=None)
```

```
sns.scatterplot(x="length_cm",  
                y="mass_g",  
                hue="extreme_l",  
                data=roach)
```




# Response values away from the regression line

```
roach["extreme_m"] = roach["mass_g"] < 1
```

```
fig = plt.figure()
sns.regplot(x="length_cm",
            y="mass_g",
            data=roach,
            ci=None)
```

```
sns.scatterplot(x="length_cm",
                y="mass_g",
                hue="extreme_l",
                style="extreme_m",
                data=roach)
```



# Leverage and influence

*Leverage* is a measure of how extreme the explanatory variable values are.

*Influence* measures how much the model would change if you left the observation out of the dataset when modeling.



# .get\_influence() and .summary\_frame()

```
mdl_roach = ols("mass_g ~ length_cm", data=roach).fit()
summary_roach = mdl_roach.get_influence().summary_frame()
roach["leverage"] = summary_roach["hat_diag"]

print(roach.head())
```

	species	mass_g	length_cm	leverage
35	Roach	40.0	12.9	0.313729
36	Roach	69.0	16.5	0.125538
37	Roach	78.0	17.5	0.093487
38	Roach	87.0	18.2	0.076283
39	Roach	120.0	18.6	0.068387

# Cook's distance

*Cook's distance* is the most common measure of influence.

```
roach["cooks_dist"] = summary_roach["cooks_d"]  
print(roach.head())
```

	species	mass_g	length_cm	leverage	cooks_dist
35	Roach	40.0	12.9	0.313729	1.074015
36	Roach	69.0	16.5	0.125538	0.010429
37	Roach	78.0	17.5	0.093487	0.000020
38	Roach	87.0	18.2	0.076283	0.001980
39	Roach	120.0	18.6	0.068387	0.006610

# Most influential roaches

```
print(roach.sort_values("cooks_dist", ascending = False))
```


	species	mass_g	length_cm	leverage	cooks_dist	
35	Roach	40.0	12.9	0.313729	1.074015	# really short roach
54	Roach	390.0	29.5	0.394740	0.365782	# really long roach
40	Roach	0.0	19.0	0.061897	0.311852	# roach with zero mass
52	Roach	290.0	24.0	0.099488	0.150064	
51	Roach	180.0	23.6	0.088391	0.061209	
..	...	...	...	...	...	
43	Roach	150.0	20.4	0.050264	0.000257	
44	Roach	145.0	20.5	0.050092	0.000256	
42	Roach	120.0	19.4	0.056815	0.000199	
47	Roach	160.0	21.1	0.050910	0.000137	
37	Roach	78.0	17.5	0.093487	0.000020	

# Removing the most influential roach

```
roach_not_short = roach[roach["length_cm"] != 12.9]
```

```
sns.regplot(x="length_cm",  
            y="mass_g",  
            data=roach,  
            ci=None,  
            line_kws={"color": "green"})
```

```
sns.regplot(x="length_cm",  
            y="mass_g",  
            data=roach_not_short,  
            ci=None,  
            line_kws={"color": "red"})
```





# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Why you need logistic regression

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# Bank churn dataset

has_churned	time_since_first_purchase	time_since_last_purchase
0	0.3993247	-0.5158691
1	-0.4297957	0.6780654
0	3.7383122	0.4082544
0	0.6032289	-0.6990435
...	...	...
<i>response</i>	<i>length of relationship</i>	<i>recency of activity</i>

<sup>1</sup> <https://www.rdocumentation.org/packages/bayesQR/topics/Churn>

# Churn vs. recency: a linear model


```
mdl_churn_vs_recency_lm = ols("has_churned ~ time_since_last_purchase",  
                               data=churn).fit()  
  
print(mdl_churn_vs_recency_lm.params)
```

```
Intercept          0.490780  
time_since_last_purchase  0.063783  
dtype: float64
```

```
intercept, slope = mdl_churn_vs_recency_lm.params
```


# Visualizing the linear model

```
sns.scatterplot(x="time_since_last_purchase",  
               y="has_churned",  
               data=churn)  
  
plt.axline(xy1=(0, intercept),  
           slope=slope)  
  
plt.show()
```



# Zooming out

```
sns.scatterplot(x="time_since_last_purchase",  
               y="has_churned",  
               data=churn)  
  
plt.axline(xy1=(0, intercept),  
           slope=slope)  
  
plt.xlim(-10, 10)  
plt.ylim(-0.2, 1.2)  
plt.show()
```



# What is logistic regression?

- Another type of generalized linear model.
- Used when the response variable is logical.
- The responses follow logistic (S-shaped) curve.

# Logistic regression using logit()

```
from statsmodels.formula.api import logit
mdl_churn_vs_recency_logit = logit("has_churned ~ time_since_last_purchase",
                                   data=churn).fit()
print(mdl_churn_vs_recency_logit.params)
```

```
Intercept          -0.035019
time_since_last_purchase    0.269215
dtype: float64
```




# Visualizing the logistic model


```
sns.regplot(x="time_since_last_purchase",  
            y="has_churned",  
            data=churn,  
            ci=None,  
            logistic=True)
```

```
plt.axline(xy1=(0, intercept),  
           slope=slope,  
           color="black")
```

```
plt.show()
```



# Zooming out



# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Predictions and odds ratios

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON




**Maarten Van den Broeck**

Content Developer at DataCamp

# The regplot() predictions

```
sns.regplot(x="time_since_last_purchase",  
            y="has_churned",  
            data=churn,  
            ci=None,  
            logistic=True)  
  
plt.show()
```




# Making predictions

```
mdl_recency = logit("has_churned ~ time_since_last_purchase",  
                    data = churn).fit()  
  
explanatory_data = pd.DataFrame(  
    {"time_since_last_purchase": np.arange(-1, 6.25, 0.25)})  
  
prediction_data = explanatory_data.assign(  
    has_churned = mdl_recency.predict(explanatory_data))
```

# Adding point predictions

```
sns.regplot(x="time_since_last_purchase",  
            y="has_churned",  
            data=churn,  
            ci=None,  
            logistic=True)  
  
sns.scatterplot(x="time_since_last_purchase",  
                y="has_churned",  
                data=prediction_data,  
                color="red")  
  
plt.show()
```



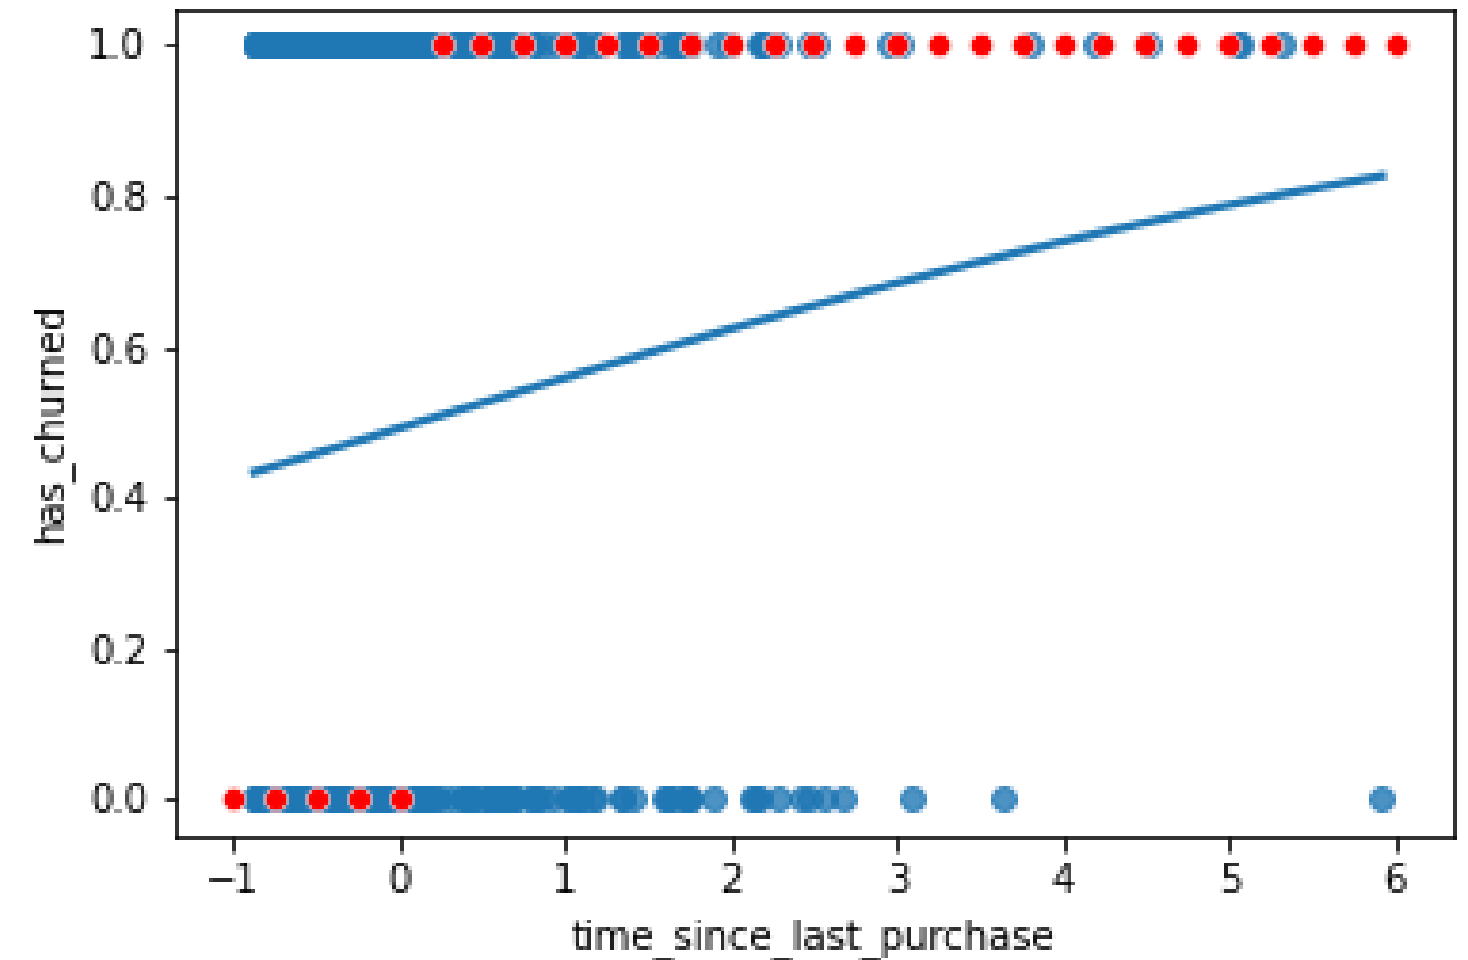
# Getting the most likely outcome

```
prediction_data = explanatory_data.assign(  
    has_churned = mdl_recency.predict(explanatory_data))  
prediction_data["most_likely_outcome"] = np.round(prediction_data["has_churned"])
```



# Visualizing most likely outcome

```
sns.regplot(x="time_since_last_purchase",  
            y="has_churned",  
            data=churn,  
            ci=None,  
            logistic=True)  
  
sns.scatterplot(x="time_since_last_purchase",  
                y="most_likely_outcome",  
                data=prediction_data,  
                color="red")  
  
plt.show()
```




# Odds ratios

*Odds ratio* is the probability of something happening divided by the probability that it doesn't.

$$\text{odds\_ratio} = \frac{\text{probability}}{(1 - \text{probability})}$$

$$\text{odds\_ratio} = \frac{0.25}{(1 - 0.25)} = \frac{1}{3}$$

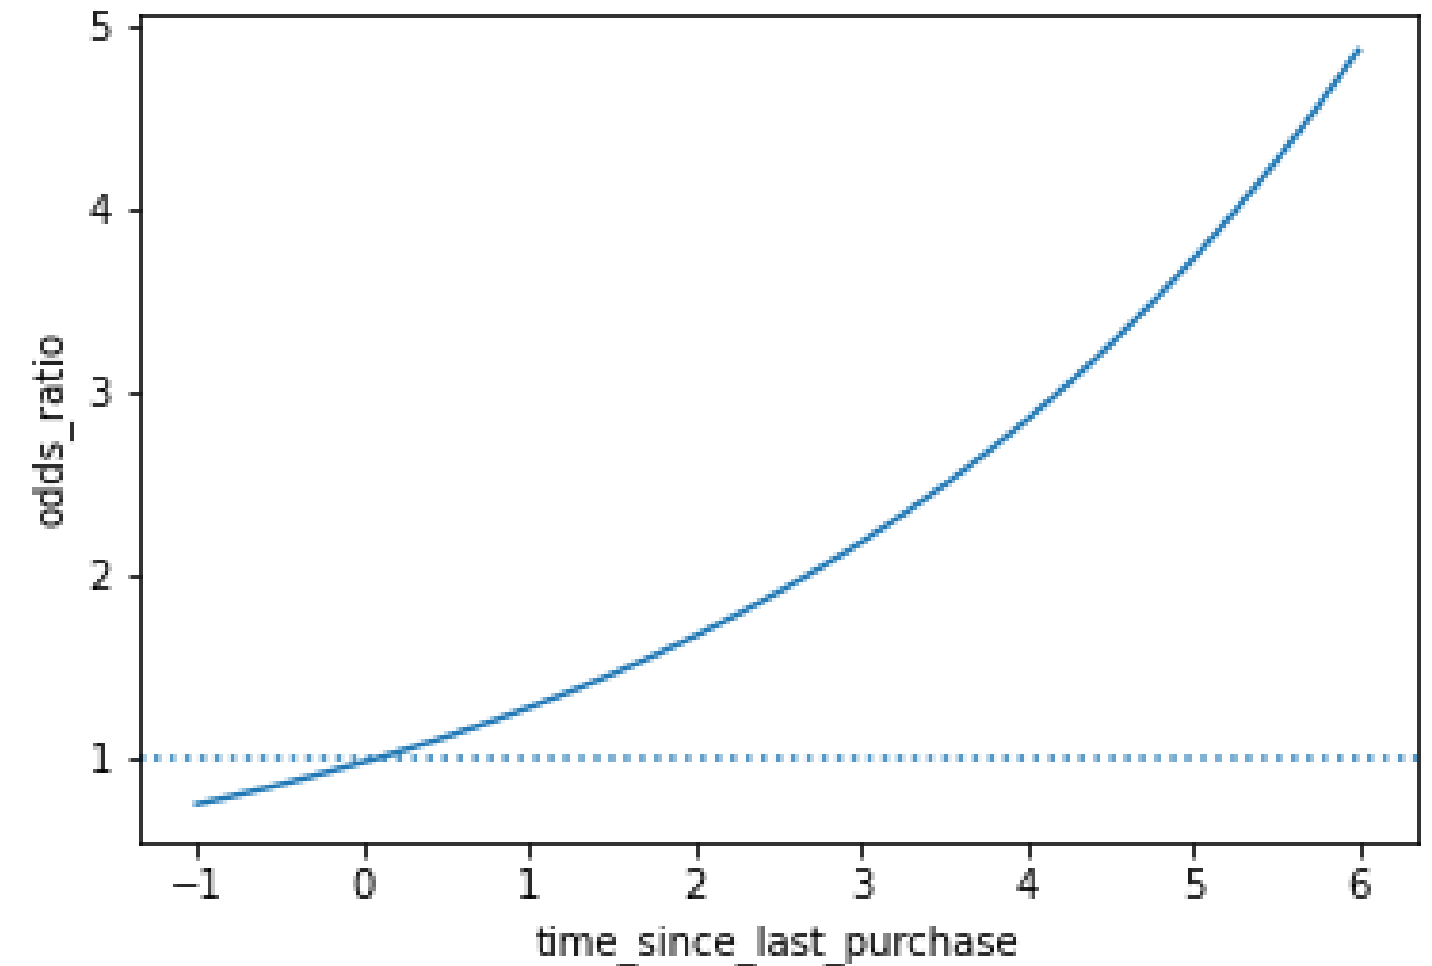


# Calculating odds ratio

```
prediction_data["odds_ratio"] = prediction_data["has_churned"] /  
                                (1 - prediction_data["has_churned"])
```

# Visualizing odds ratio

```
sns.lineplot(x="time_since_last_purchase",  
             y="odds_ratio",  
             data=prediction_data)  
  
plt.axhline(y=1,  
            linestyle="dotted")  
  
plt.show()
```




# Visualizing log odds ratio

```
sns.lineplot(x="time_since_last_purchase",  
             y="odds_ratio",  
             data=prediction_data)
```

```
plt.axhline(y=1,  
            linestyle="dotted")
```

```
plt.yscale("log")
```

```
plt.show()
```



# Calculating log odds ratio

```
prediction_data["log_odds_ratio"] = np.log(prediction_data["odds_ratio"])
```

# All predictions together

time_since_last_prchs	has_churned	most_likely_rspns	odds_ratio	log_odds_ratio
0	0.491	0	0.966	-0.035
2	0.623	1	1.654	0.503
4	0.739	1	2.834	1.042
6	0.829	1	4.856	1.580
...	...	...	...	...

# Comparing scales

Scale	Are values easy to interpret?	Are changes easy to interpret?	Is precise?
Probability	✓	✗	✓
Most likely outcome	✓✓	✓	✗
Odds ratio	✓	✗	✓
Log odds ratio	✗	✓	✓



# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Quantifying logistic regression fit

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# The four outcomes

	predicted false	predicted true
actual false	correct	false positive
actual true	false negative	correct

# Confusion matrix: counts of outcomes

```
actual_response = churn["has_churned"]
```

```
predicted_response = np.round mdl_recency.predict())
```

```
outcomes = pd.DataFrame({"actual_response": actual_response,  
                          "predicted_response": predicted_response})
```

```
print(outcomes.value_counts(sort=False))
```

actual_response	predicted_response	
0	0.0	141
	1.0	59
1	0.0	111
	1.0	89

# Visualizing the confusion matrix

```
conf_matrix = mdl_recency.pred_table()

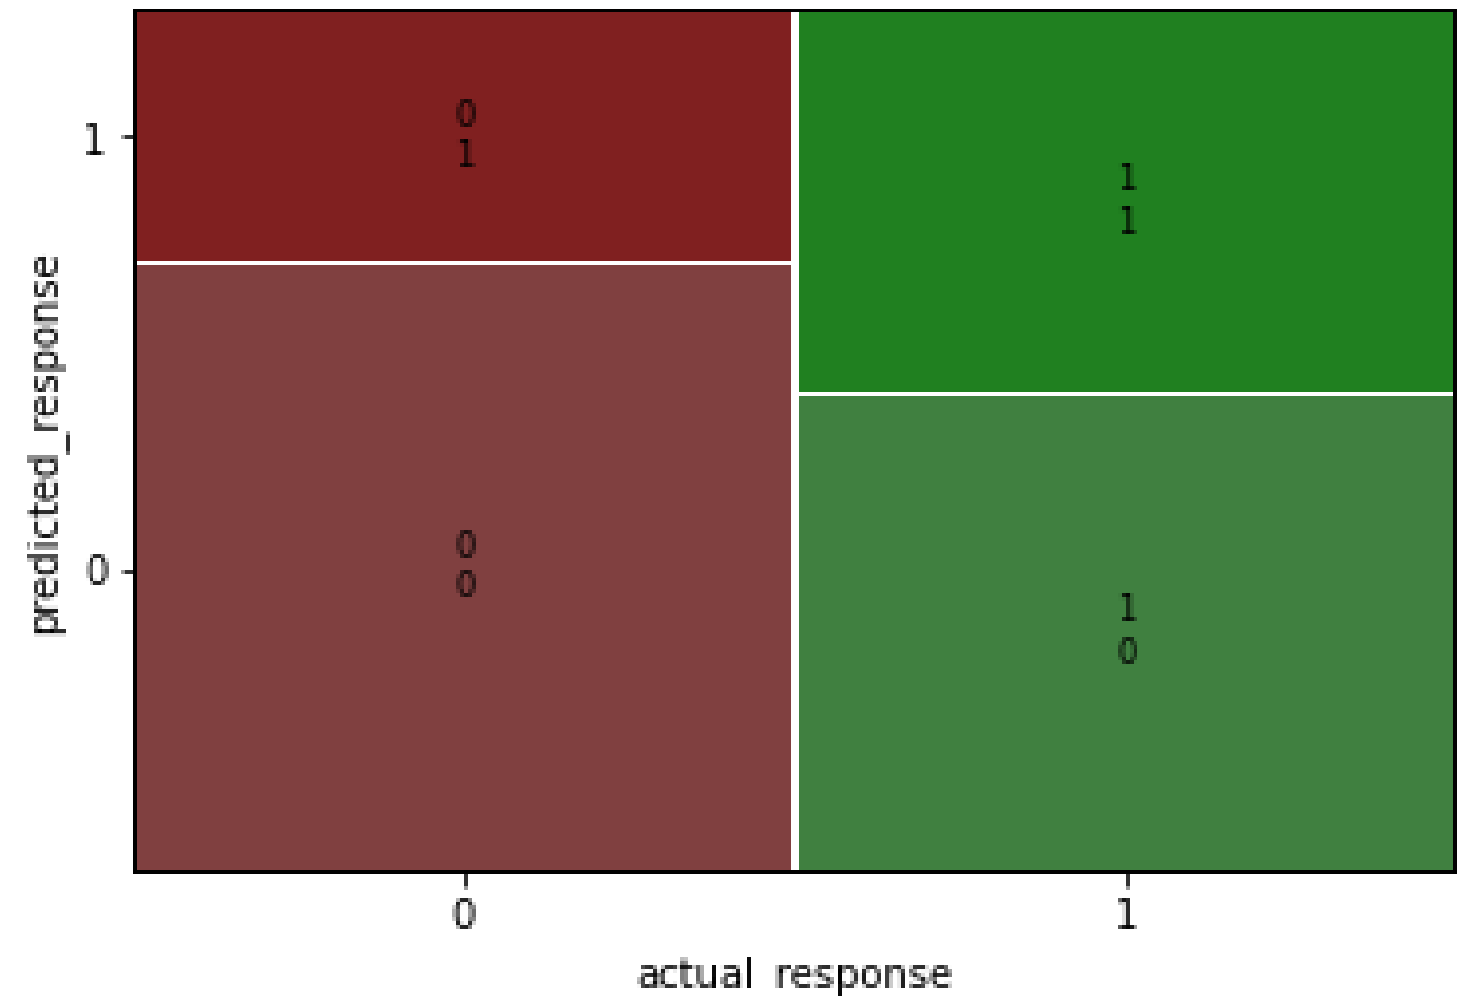
print(conf_matrix)
```

```
[[141.    59.]
 [111.    89.]]
```

true negative	false positive
false negative	true positive

```
from statsmodels.graphics.mosaicplot
import mosaic

mosaic(conf_matrix)
```



# Accuracy

*Accuracy* is the proportion of correct predictions.

$$\text{accuracy} = \frac{TN + TP}{TN + FN + FP + TP}$$

```
[[141.,  59.],  
 [111.,  89.]]
```

```
TN = conf_matrix[0,0]  
TP = conf_matrix[1,1]  
FN = conf_matrix[1,0]  
FP = conf_matrix[0,1]
```

```
acc = (TN + TP) / (TN + TP + FN + FP)  
print(acc)
```

```
0.575
```

# Sensitivity

*Sensitivity* is the proportion of true positives.

$$\text{sensitivity} = \frac{TP}{FN + TP}$$

```
[[141.,  59.],  
 [111.,  89.]]
```

```
TN = conf_matrix[0,0]  
TP = conf_matrix[1,1]  
FN = conf_matrix[1,0]  
FP = conf_matrix[0,1]
```

```
sens = TP / (FN + TP)  
print(sens)
```

```
0.445
```

# Specificity

*Specificity* is the proportion of true negatives.

$$\text{specificity} = \frac{TN}{TN + FP}$$

```
[[141.,  59.],  
 [111.,  89.]]
```

```
TN = conf_matrix[0,0]  
TP = conf_matrix[1,1]  
FN = conf_matrix[1,0]  
FP = conf_matrix[0,1]
```

```
spec = TN / (TN + FP)  
print(spec)
```

```
0.705
```



# Let's practice!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON

# Congratulations

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON



**Maarten Van den Broeck**

Content Developer at DataCamp

# You learned things

## Chapter 1

- Fit a simple linear regression
- Interpret coefficients

## Chapter 3

- Quantifying model fit
- Outlier, leverage, and influence

## Chapter 2

- Make predictions
- Regression to the mean
- Transforming variables

## Chapter 4

- Fit a simple logistic regression
- Make predictions
- Get performance from confusion matrix

# Multiple explanatory variables

Intermediate Regression with statsmodels in Python

# Unlocking advanced skills

- Generalized Linear Models in Python
- Introduction to Predictive Analytics in Python
- Linear Classifiers in Python

# Happy learning!

INTRODUCTION TO REGRESSION WITH STATSMODELS IN PYTHON